

Main Memory: Part 1

COMP 3659 – Operating Systems

Monday, November 8, 2021

Text: Chapter 9 (9.1-9.2)

Announcements and Course Administration

- Programming project 2:
 - Concept memos have been received
 - Except group feedback early this week – watch your email

Overview

- Last lecture:
 - Deadlock: liveness; system model; deadlock in multithreaded applications; characterization; overview of methods for handling (text sections 6.8, 8.1-8.4)
- Starting this week:
 - A two-week unit on memory management: textbook chapters 9-10
- Today's focuses:
 - Background 9.1
 - Contiguous memory allocation 9.2

Lab Debrief

- Discuss: What key things did you learn in the lab?
 - Summary of key learnings
 - About phenomena, interpretations, concepts, emerging themes, ...
 - Insights? Questions? New goals?
 - Challenges and successes?
 - Other?


Last Lecture Recap:

- Deadlock
 - A main category of liveness problem (note also livelock)
 - >1 thread waiting for an event that can only be signaled by one of the waiting threads
- Often involves the incorrect introduction of thread synchronization
 - E.g., using semaphores
 - Attempting to solve a race condition can introduce the potential for deadlock!
- Can be visualized using resource allocation graph
 - No cycle → no deadlock; cycle → potential for deadlock (not always)
- E.g., the dining philosophers problem and one potential solution
- Deadlock prevention, avoidance, detection, and recovery
 - kernel can also ignore the possibility of deadlock



Background

- ❑ Program must be brought (from disk) into memory and placed within a process for it to be run
- ❑ Main memory and registers are only storage CPU can access directly
- ❑ Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- ❑ Register access in one CPU clock (or less)
- ❑ Main memory can take many cycles, causing a **stall**
- ❑ **Cache** sits between main memory and CPU registers
- ❑ Protection of memory required to ensure correct operation




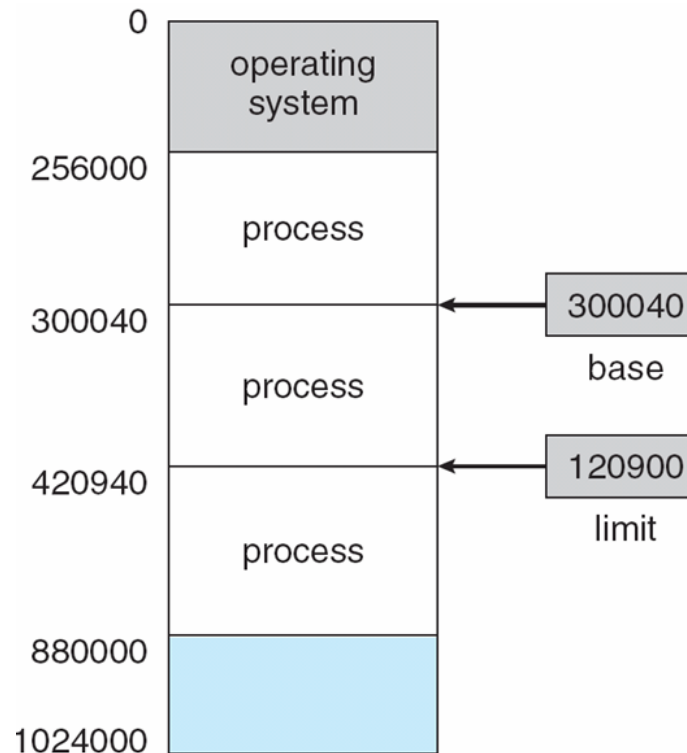
The Simplest Model of Memory Protection

- Multiple process (and the kernel) all co-exist in physical memory
- The operating system allocates non-overlapping regions of physical memory to each
 - Can consider each region to be that process's "logical address space"
- Implemented via:
 - Base register
 - Limit register
- Basically, each process executes in an "address space jail"



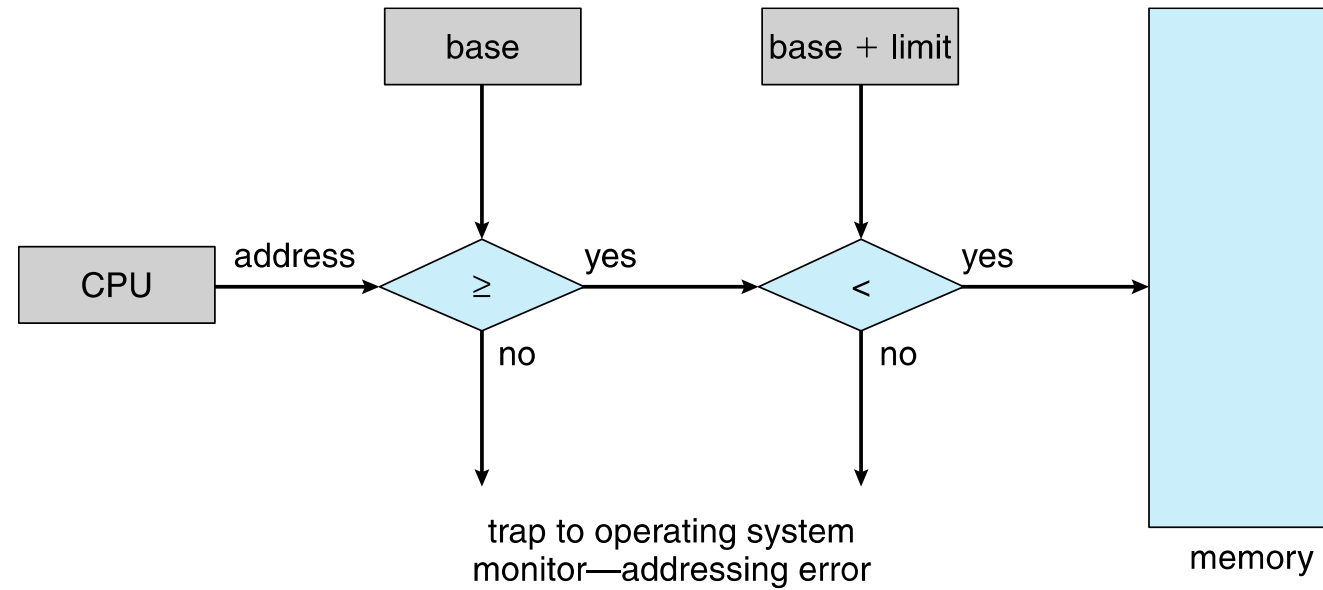
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user






Hardware Address Protection





Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
 - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
 - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses **bind** to relocatable addresses
 - ▶ i.e. "14 bytes from beginning of this module"
 - Linker or loader will bind relocatable addresses to absolute addresses
 - ▶ i.e. 74014
 - Each binding maps one address space to another





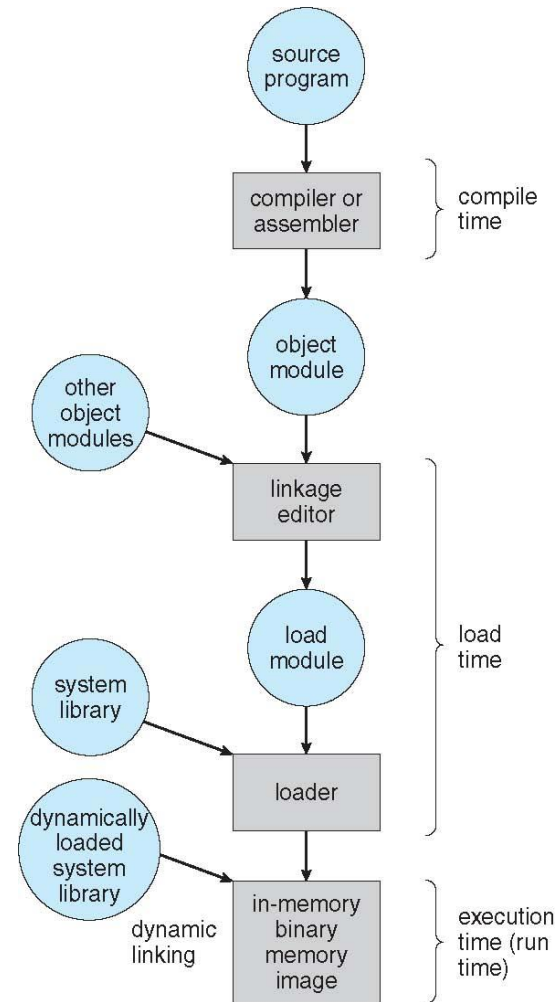
Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address maps (e.g., base and limit registers)





Multistep Processing of a User Program





Logical vs. Physical Address Space

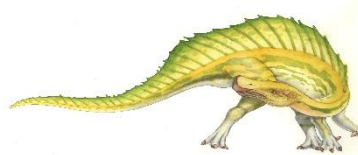
- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





Memory-Management Unit (MMU)



- ❑ Hardware device that at run time maps virtual to physical address
- ❑ Many methods possible, covered in the rest of this chapter
- ❑ To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - ❑ Base register now called **relocation register**
 - ❑ MS-DOS on Intel 80x86 used 4 relocation registers
- ❑ The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - ❑ Execution-time binding occurs when reference is made to location in memory
 - ❑ Logical address bound to physical addresses





Dynamic relocation using a relocation register

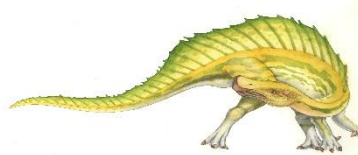
- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
 - Implemented through program design
 - OS can help by providing libraries to implement dynamic loading





Dynamic Linking


- ❑ **Static linking** – system libraries and program code combined by the loader into the binary program image
- ❑ Dynamic linking –linking postponed until execution time
- ❑ Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- ❑ Stub replaces itself with the address of the routine, and executes the routine
- ❑ Operating system checks if routine is in processes' memory address
 - ❑ If not in address space, add to address space
- ❑ Dynamic linking is particularly useful for libraries
- ❑ System also known as **shared libraries**
- ❑ Consider applicability to patching system libraries
 - ❑ Versioning may be needed





Contiguous Allocation

- ❑ Main memory must support both OS and user processes
- ❑ Limited resource, must allocate efficiently
- ❑ Contiguous allocation is one early method
- ❑ Main memory usually into two **partitions**:
 - ❑ Resident operating system, usually held in low memory with interrupt vector
 - ❑ User processes then held in high memory
 - ❑ Each process contained in single contiguous section of memory





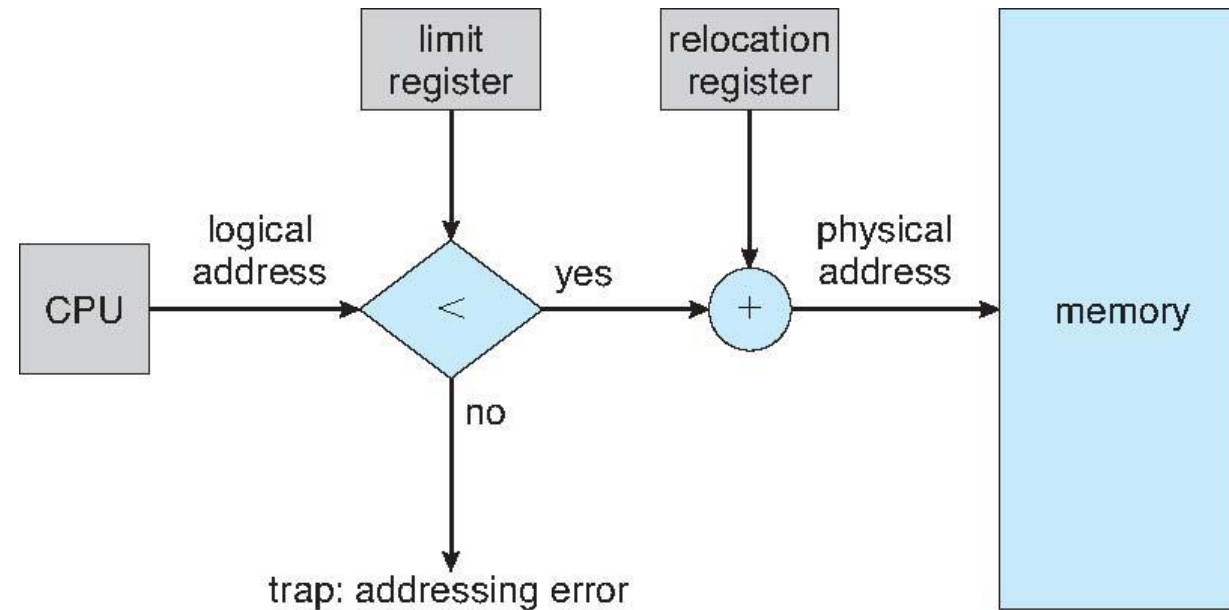
Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Can then allow actions such as kernel code being **transient** and kernel changing size





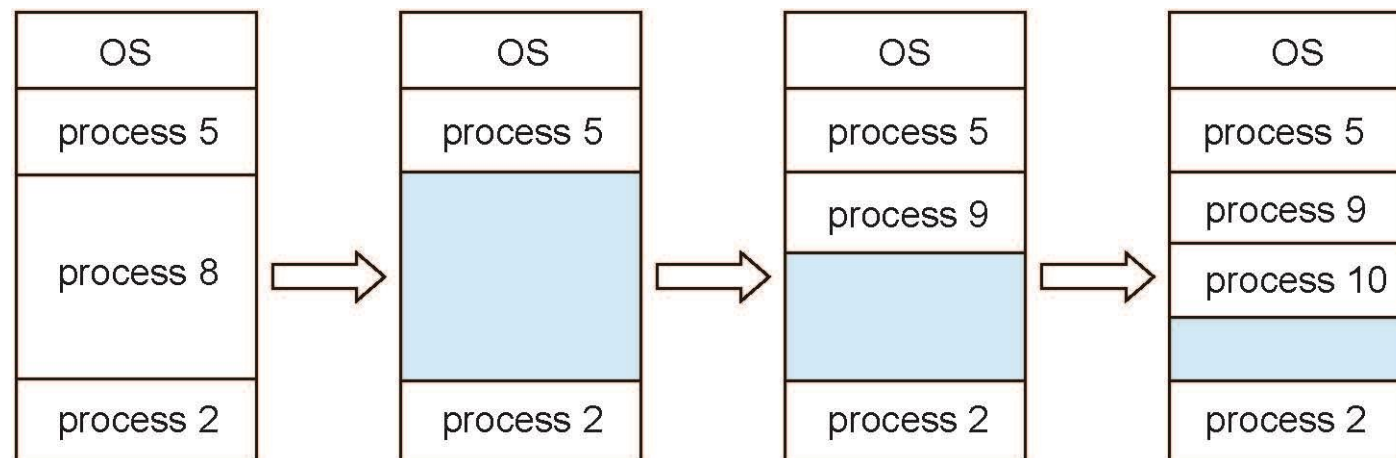
Hardware Support for Relocation and Limit Registers





Multiple-partition allocation

- ❑ Multiple-partition allocation
 - ❑ Degree of multiprogramming limited by number of partitions
 - ❑ **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - ❑ **Hole** – block of available memory; holes of various size are scattered throughout memory
 - ❑ When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - ❑ Process exiting frees its partition, adjacent free partitions combined
 - ❑ Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)






Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- **First-fit**: Allocate the **first** hole that is big enough
- **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit**: Allocate the **largest** hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





Fragmentation


- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable -> **50-percent rule**





Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - I/O problem
 - ▶ Latch job in memory while it is involved in I/O
 - ▶ Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems



Conclusion

- A modern kernel must allocate space to ≥ 1 processes (and itself)
 - While providing memory protection
 - Raises allocation strategy and fragmentation issues
- One “simple” solution: contiguous allocation with base & limit registers
- A program’s instructions & data must be bound to the (logical) address range provided by the kernel
 - Can only be done at compile time if absolute (logical) addresses known a priori
 - Must be done at load and/or execute time otherwise; code must be relocatable
- A hardware MMU supports relocation
 - And better-distinguishes the processes *logical address space* from physical addresses
- An elegant upcoming solution: paging
- Next:
 - Paging (and related topics) 9.3-9.4, 10.1-10.2

Recommended Post-class Studying

- Review these slides and your lecture notes
 - Ideally within 24-48 hours
 - Then review text sections 9.1-9.2 as necessary
- As always, please let me know if you have any questions